

Nimisha Majgawali

+91 8962763616 | majgawalinimisha@gmail.com | [LinkedIn](#) | [GitHub](#)

SUMMARY

Final-year B.Tech Computer Science student with multiple production-deployed AI/ML systems spanning multi-agent orchestration, RAG pipelines, NLP scoring, and LLM API integration. Built and shipped end-to-end applications in Python (Flask, FastAPI) and React across real internship environments. Comfortable applying statistical and ML techniques to structured datasets, designing REST APIs, and working across the full stack from data pipeline to deployed interface. Strong interest in Generative AI and how emerging techniques translate into working software.

SKILLS

AI & ML: Scikit-learn, KNN, TF-IDF, Cosine Similarity, NLP, Regression, Clustering, LangGraph, CrewAI, Prompt Engineering, Sentence Transformers, FAISS

LLM Integration: Groq LLaMA 3.1, Gemini 1.5 Flash, Azure OpenAI, RAG pipelines, LLM API calls

Languages: Python, JavaScript (ES6+), TypeScript, SQL, Bash, Java (familiar)

Frameworks & APIs: FastAPI, Flask, Node.js, Express.js, React.js, RESTful API Design, WebSockets

Databases: MySQL, MongoDB, PostgreSQL, Neo4j, FAISS

Cloud & DevOps: AWS (EC2, S3, ECS, IAM), Docker, GitHub, Vercel, Render

Frontend: React.js, React 19, TypeScript, Tailwind CSS, HTML5, CSS3, Figma, Storybook

Tools: Git, Postman, Google Colab, Agile/Scrum, Code Reviews, Technical Documentation

EXPERIENCE

Front-End Developer Intern | MediKloud (Healthcare SaaS)

Apr 2026 – May 2026

- Refactored a fragmented production React/TypeScript codebase into a component-driven design system, establishing a centralised Tailwind CSS token architecture as a single source of truth across the platform.
- Structured reusable UI components with explicit prop contracts documented in Storybook, reducing onboarding effort for future contributors and improving component isolation.
- Collaborated with designers and engineers to validate component behaviour against acceptance criteria before each release, contributing to quality-gating on a regulated healthcare software platform.
- Eliminated all hardcoded styling values across a large shared codebase, ensuring token changes propagate uniformly to every component with zero inline style overrides.

Software Engineering Intern | Ethnus

Jan 2025 – May 2025

- Delivered 4 production modules — authentication, real-time messaging, document sharing, and role-based access control — across 3 agile sprints on a live platform serving 15+ users, with zero regressions per release.
- Contributed to backend middleware and API configuration across 8 endpoints within a 4-person cross-functional team, enforcing admin and user permission hierarchies through RBAC conditional logic.
- Integrated WebSocket-based real-time collaborative features enabling synchronous editing across 6 concurrent users, coordinating frontend and backend service integration.
- Provisioned AWS EC2 instances with security group rules, IAM policies, and network isolation; wrote Bash scripts for environment setup and log aggregation.
- Reviewed approximately 4 pull requests per sprint, flagging logic gaps and edge cases before each production release.

PROJECTS

AgentGraph-Intel | *Python, LangGraph, Neo4j, FastAPI, React.js, AWS ECS* | [GitHub](#)

Jan 2026 – Feb 2026

- Architected a multi-agent AI research system in LangGraph with stateful orchestration across 3 specialised agents (Researcher, Knowledge Graph Builder, Analyst), driven by real-time intent classification and conditional routing between agents.
- Knowledge Graph Builder agent extracts entities and relationships via LLM-powered NLP, persisting to Neo4j via Cypher — enabling multi-hop reasoning across connected data that pure vector search cannot resolve.
- Deployed the full FastAPI + React stack to AWS ECS via Docker with CloudWatch monitoring and UUID-based multi-tenant session isolation; automated build-test-deploy via GitHub Actions CI/CD pipeline.

DocuMind AI | *Python, FastAPI, FAISS, BM25, Sentence Transformers, Groq LLaMA 3.1* | [GitHub](#)

Nov 2025 – Dec 2025

- Built a production RAG document intelligence platform where users upload PDF, DOCX, and TXT files and receive source-cited, confidence-scored answers via a REST API backed by a full retrieval pipeline.
- Implemented a 4-stage retrieval pipeline — semantic chunking (NLTK), hybrid FAISS + BM25 with 60/40 score fusion, cross-encoder re-ranking, and query decomposition — substantially outperforming single-vector baselines.
- Added page-level source citation via chunk-to-page mapping and a self-evaluation confidence score (1–10) per answer; engineered session-based conversation memory with UUID multi-user isolation.

SkillRadar | *Python, Flask, Scikit-learn, TF-IDF, NLP, Groq LLaMA 3.1, Docker* | [GitHub](#)

Feb 2026 – Mar 2026

- Built an NLP pipeline using TF-IDF vectorisation and cosine similarity to score resume–JD alignment, surfacing matched and missing skills separately rather than returning a raw keyword overlap count.
- Cross-referenced a 200+ tech-skills dictionary against both documents to generate actionable gap output; called Groq LLaMA 3.1 API with missing-skill context to return 3–5 targeted improvement suggestions per scan.

- Containerised the full application with Docker and exposed all functionality through a REST API, making the tool integrable as an external service rather than a standalone script.

CareCrypt | *Python, Flask, MySQL, Fernet AES-128, Groq LLaMA 3.1, Docker* | [GitHub](#) Jan 2026 – Feb 2026

- Designed a REST API backend across 13 endpoints and a 5-table normalised MySQL schema, applying strict separation between data, service, and routing layers — keeping the codebase modular and independently testable.
- Integrated Groq LLaMA 3.1 to generate plain-English prescription summaries from structured clinical records; applied field-level AES-128 encryption so plaintext patient data never persists to disk.
- Applied five independent security controls — rate limiting, bcrypt hashing, CSRF protection, parameterised queries, and time-limited tokens — sustaining sub-200ms response times across 10 beta users on a query-optimised schema.

NutriGenie | *Python, FastAPI, Scikit-learn, KNN, MongoDB, Docker, Hugging Face* | [GitHub](#) Dec 2024 – Feb 2025

- Architected a FastAPI backend running a KNN + cosine similarity model over a 522K+ recipe dataset, returning ranked personalised meal plans as structured JSON within 300ms average response time.
- Designed a 6-parameter user profiling pipeline computing BMI, BMR, and caloric targets across 4 weight plans; separated Streamlit frontend and FastAPI backend into independently deployable services.

EDUCATION

VIT Bhopal University

B.Tech – Computer Science (Minor: Health Informatics); GPA: 8.63/10

Bhopal, India

2022 – 2026

DPS Durg

Class XII CBSE: 94% | Class X CBSE: 91%

Chhattisgarh, India

2018 – 2020

RELEVANT COURSEWORK

Core CS: Data Structures & Algorithms, Operating Systems, Computer Networks, Database Management Systems, Object-Oriented Design

AI & Data: Applied Machine Learning, Natural Language Processing, AI model development, Statistical Methods

Health Informatics Minor: Healthcare data standards (HIPAA, GDPR), Electronic Health Records, Clinical Decision Support Systems, IoT in Healthcare

CERTIFICATIONS

[Applied Machine Learning in Python](#) – University of Michigan

[Cloud Architecting](#) – AWS Academy Graduate

[Cloud Foundations](#) – AWS Academy Graduate

[AWS Cloud Fundamentals](#) – Ethnus Training Program (2025)